

Source Code

September 28, 2018

```
In [1]: #code initialization
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn import preprocessing
import seaborn as sns
from sklearn.cluster import KMeans
from sklearn.metrics import mean_squared_error
from math import sqrt

pd.set_option('display.float_format', lambda x: '%.5f' % x)
```

```
In [2]: #read data
data = pd.read_csv("creditcard.csv")
#data = pd.read_csv("creditcard.csv", low_memory=False)
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284909 entries, 0 to 284908
Data columns (total 31 columns):
Time          284909 non-null int64
V1            284909 non-null object
V2            284909 non-null object
V3            284909 non-null object
V4            284909 non-null object
V5            284909 non-null object
V6            284909 non-null object
V7            284909 non-null float64
V8            284909 non-null float64
V9            284909 non-null float64
V10           284909 non-null float64
V11           284909 non-null float64
V12           284909 non-null float64
V13           284909 non-null float64
V14           284909 non-null float64
V15           284909 non-null float64
V16           284909 non-null float64
V17           284909 non-null float64
```

```

V18      284909 non-null float64
V19      284909 non-null float64
V20      284909 non-null float64
V21      284909 non-null float64
V22      284909 non-null float64
V23      284909 non-null float64
V24      284909 non-null object
V25      284909 non-null object
V26      284909 non-null float64
V27      284909 non-null float64
V28      284909 non-null float64
Amount   284909 non-null float64
Class    284909 non-null int64
dtypes: float64(21), int64(2), object(8)
memory usage: 67.4+ MB

```

```

/home/vader13/anaconda3/lib/python3.6/site-packages/IPython/core/interactiveshell.py:2728: DtypeWarning:
  interactivity=interactivity, compiler=compiler, result=result)

```

```

In [3]: #replace objects with NaN
        column = data.columns[data.dtypes.eq(object)]

        for col in column:
            data[col] = pd.to_numeric(data[col], errors='coerce')

        data.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284909 entries, 0 to 284908
Data columns (total 31 columns):
Time      284909 non-null int64
V1        284907 non-null float64
V2        284908 non-null float64
V3        284908 non-null float64
V4        284908 non-null float64
V5        284907 non-null float64
V6        284908 non-null float64
V7        284909 non-null float64
V8        284909 non-null float64
V9        284909 non-null float64
V10       284909 non-null float64
V11       284909 non-null float64
V12       284909 non-null float64
V13       284909 non-null float64
V14       284909 non-null float64
V15       284909 non-null float64

```

```

V16      284909 non-null float64
V17      284909 non-null float64
V18      284909 non-null float64
V19      284909 non-null float64
V20      284909 non-null float64
V21      284909 non-null float64
V22      284909 non-null float64
V23      284909 non-null float64
V24      284894 non-null float64
V25      284903 non-null float64
V26      284909 non-null float64
V27      284909 non-null float64
V28      284909 non-null float64
Amount   284909 non-null float64
Class    284909 non-null int64
dtypes: float64(29), int64(2)
memory usage: 67.4 MB

```

```

In [4]: #Drop NaN values
        data.dropna(inplace = True)
        #data.fillna(data.mean(), inplace=True)

```

```

data.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 284880 entries, 0 to 284907
Data columns (total 31 columns):
Time      284880 non-null int64
V1        284880 non-null float64
V2        284880 non-null float64
V3        284880 non-null float64
V4        284880 non-null float64
V5        284880 non-null float64
V6        284880 non-null float64
V7        284880 non-null float64
V8        284880 non-null float64
V9        284880 non-null float64
V10       284880 non-null float64
V11       284880 non-null float64
V12       284880 non-null float64
V13       284880 non-null float64
V14       284880 non-null float64
V15       284880 non-null float64
V16       284880 non-null float64
V17       284880 non-null float64
V18       284880 non-null float64
V19       284880 non-null float64

```

```

V20      284880 non-null float64
V21      284880 non-null float64
V22      284880 non-null float64
V23      284880 non-null float64
V24      284880 non-null float64
V25      284880 non-null float64
V26      284880 non-null float64
V27      284880 non-null float64
V28      284880 non-null float64
Amount   284880 non-null float64
Class    284880 non-null int64
dtypes: float64(29), int64(2)
memory usage: 69.6 MB

```

```

In [5]: #Drop Duplicate values
        data = data.drop_duplicates()
        data.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 283726 entries, 0 to 284807
Data columns (total 31 columns):
Time      283726 non-null int64
V1        283726 non-null float64
V2        283726 non-null float64
V3        283726 non-null float64
V4        283726 non-null float64
V5        283726 non-null float64
V6        283726 non-null float64
V7        283726 non-null float64
V8        283726 non-null float64
V9        283726 non-null float64
V10       283726 non-null float64
V11       283726 non-null float64
V12       283726 non-null float64
V13       283726 non-null float64
V14       283726 non-null float64
V15       283726 non-null float64
V16       283726 non-null float64
V17       283726 non-null float64
V18       283726 non-null float64
V19       283726 non-null float64
V20       283726 non-null float64
V21       283726 non-null float64
V22       283726 non-null float64
V23       283726 non-null float64
V24       283726 non-null float64
V25       283726 non-null float64

```

```
V26      283726 non-null float64
V27      283726 non-null float64
V28      283726 non-null float64
Amount   283726 non-null float64
Class    283726 non-null int64
dtypes: float64(29), int64(2)
memory usage: 69.3 MB
```

```
In [6]: #Amount and Time cannot be negative by domain knowledge
```

```
data = data[data.Amount>=0]
```

```
data = data[data.Time>=0]
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 283726 entries, 0 to 284807
Data columns (total 31 columns):
Time      283726 non-null int64
V1        283726 non-null float64
V2        283726 non-null float64
V3        283726 non-null float64
V4        283726 non-null float64
V5        283726 non-null float64
V6        283726 non-null float64
V7        283726 non-null float64
V8        283726 non-null float64
V9        283726 non-null float64
V10       283726 non-null float64
V11       283726 non-null float64
V12       283726 non-null float64
V13       283726 non-null float64
V14       283726 non-null float64
V15       283726 non-null float64
V16       283726 non-null float64
V17       283726 non-null float64
V18       283726 non-null float64
V19       283726 non-null float64
V20       283726 non-null float64
V21       283726 non-null float64
V22       283726 non-null float64
V23       283726 non-null float64
V24       283726 non-null float64
V25       283726 non-null float64
V26       283726 non-null float64
V27       283726 non-null float64
V28       283726 non-null float64
Amount    283726 non-null float64
```

```

Class      283726 non-null int64
dtypes: float64(29), int64(2)
memory usage: 69.3 MB

```


```

In [7]: #Heatmap for visualization
f, ax = plt.subplots(figsize=(12,8))
correlate = data.corr()

sns.heatmap(correlate, mask=np.zeros_like(correlate, dtype=np.bool), \
            cmap=sns.diverging_palette(220, 10, as_cmap=True), \
            square=True, ax=ax)

plt.title('Fig1. Heat Map')
plt.xlabel('Attributes')
plt.ylabel('Attributes')
plt.show()

```



```
In [8]: #Find magnitude of correlation in sorted order  
abs(correlate.Class).sort_values()
```

```
Out[8]: V25      0.00320  
V15      0.00330  
V13      0.00390  
V26      0.00426  
V22      0.00489  
Amount   0.00578  
V23      0.00633  
V24      0.00721  
V28      0.00968  
Time     0.01236  
V20      0.02149  
V27      0.02189  
V21      0.02636  
V8       0.03307  
V19      0.03363  
V6       0.04392  
V2       0.08462  
V5       0.08781  
V9       0.09402  
V1       0.09449  
V18      0.10534  
V4       0.12933  
V11      0.14907  
V7       0.17235  
V3       0.18232  
V16      0.18719  
V10      0.20697  
V12      0.25071  
V14      0.29338  
V17      0.31350  
Class    1.00000  
Name: Class, dtype: float64
```

```
In [9]: #Create two working copies of data  
traindata1 = data.copy()  
traindata2 = data.copy()
```

```
In [10]: #traindata1 has all attribites except Class  
traindata1 = traindata1.drop(labels=['Class'],axis=1)  
traindata1.shape
```

```
Out[10]: (283726, 30)
```

```
In [11]: #Performing Min_Max Normalization on data  
min_max_scaler = preprocessing.MinMaxScaler()  
np_scaled = min_max_scaler.fit_transform(data)
```

```
dataN = pd.DataFrame(np_scaled)
dataN.head()
```

```
Out[11]:
```

	0	1	2	3	4	5	6	7	8	\
0	0.00000	0.93519	0.76649	0.88136	0.31302	0.76344	0.26767	0.26682	0.78644	
1	0.00000	0.97854	0.77007	0.84030	0.27180	0.76612	0.26219	0.26488	0.78630	
2	0.00001	0.93522	0.75312	0.86814	0.26877	0.76233	0.28112	0.27018	0.78804	
3	0.00001	0.94188	0.76530	0.86848	0.21366	0.76565	0.27556	0.26680	0.78943	
4	0.00001	0.93862	0.77652	0.86425	0.26980	0.76298	0.26398	0.26897	0.78248	
	9	...	21	22	23	24	25	26	27	\
0	0.47531	...	0.56118	0.52299	0.66379	0.39125	0.58512	0.39456	0.41898	
1	0.45398	...	0.55784	0.48024	0.66694	0.33644	0.58729	0.44601	0.41635	
2	0.41060	...	0.56548	0.54603	0.67894	0.28935	0.55952	0.40273	0.41549	
3	0.41500	...	0.55973	0.51028	0.66261	0.22383	0.61425	0.38920	0.41767	
4	0.49095	...	0.56133	0.54727	0.66339	0.40127	0.56634	0.50750	0.42056	
	28	29	30							
0	0.31270	0.00582	0.00000							
1	0.31342	0.00010	0.00000							
2	0.31191	0.01474	0.00000							
3	0.31437	0.00481	0.00000							
4	0.31749	0.00272	0.00000							

[5 rows x 31 columns]

```
In [12]: #Performing Min_Max Normalization on traindata1
min_max_scaler1 = preprocessing.MinMaxScaler()
np_scaled1 = min_max_scaler1.fit_transform(traindata1)
traindata1N = pd.DataFrame(np_scaled1)
traindata1N.head()
```

```
Out[12]:
```

	0	1	2	3	4	5	6	7	8	\
0	0.00000	0.93519	0.76649	0.88136	0.31302	0.76344	0.26767	0.26682	0.78644	
1	0.00000	0.97854	0.77007	0.84030	0.27180	0.76612	0.26219	0.26488	0.78630	
2	0.00001	0.93522	0.75312	0.86814	0.26877	0.76233	0.28112	0.27018	0.78804	
3	0.00001	0.94188	0.76530	0.86848	0.21366	0.76565	0.27556	0.26680	0.78943	
4	0.00001	0.93862	0.77652	0.86425	0.26980	0.76298	0.26398	0.26897	0.78248	
	9	...	20	21	22	23	24	25	26	\
0	0.47531	...	0.58294	0.56118	0.52299	0.66379	0.39125	0.58512	0.39456	
1	0.45398	...	0.57953	0.55784	0.48024	0.66694	0.33644	0.58729	0.44601	
2	0.41060	...	0.58586	0.56548	0.54603	0.67894	0.28935	0.55952	0.40273	
3	0.41500	...	0.57805	0.55973	0.51028	0.66261	0.22383	0.61425	0.38920	
4	0.49095	...	0.58462	0.56133	0.54727	0.66339	0.40127	0.56634	0.50750	
	27	28	29							
0	0.41898	0.31270	0.00582							


```

1 0.41635 0.31342 0.00010
2 0.41549 0.31191 0.01474
3 0.41767 0.31437 0.00481
4 0.42056 0.31749 0.00272

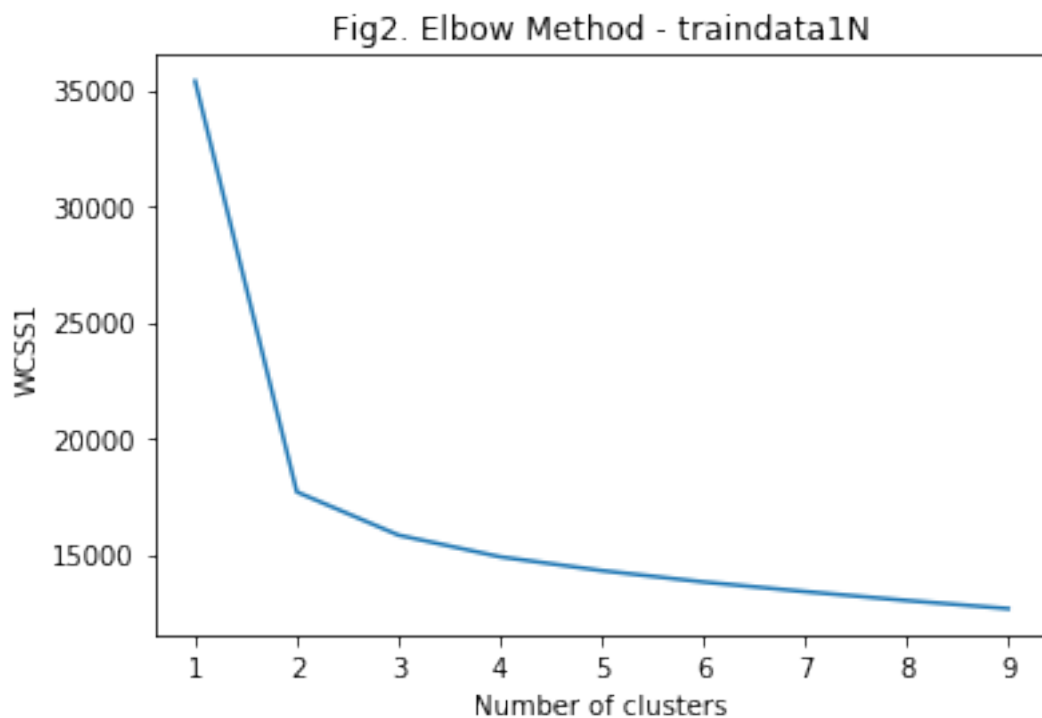
```

[5 rows x 30 columns]

```

In [13]: #Elbow method to find k for traindata1N
wcsc1 = []
for i in range(1,10):
    kmeans1 = KMeans(n_clusters = i,init = 'k-means++',random_state = 0)
    kmeans1.fit(traindata1N)
    wcsc1.append(kmeans1.inertia_)
plt.plot(range(1,10),wcsc1)
plt.title('Fig2. Elbow Method - traindata1N')
plt.xlabel('Number of clusters')
plt.ylabel('WCSS1')
plt.show()

```



```

In [14]: #K-means clustering for traindata1N with k=2
kmeans1 = KMeans(n_clusters = 2,init = 'k-means++',random_state=0)
clusters1 = kmeans1.fit_predict(traindata1N)
clusters1

```

```
Out[14]: array([1, 1, 1, ..., 0, 0, 0], dtype=int32)
```

```
In [15]: #Accuracy of Classification.  
#Consider cluster labeled by '0' to be fraudulent  
#and that by '1' to be Non-fraudulent  
count1 = 0  
for i in range(283726):  
    if dataN.values[:,30][i]==0 and clusters1[i]==1:  
        count1+=1  
    if dataN.values[:,30][i]==1 and clusters1[i]==0:  
        count1+=1  
Accuracy1 = (count1/283726)*100  
Accuracy1
```

```
Out[15]: 53.65493469051127
```

```
In [16]: #Precision of Fraudulent data classification  
#Calculated as Precision = Number of fraudulent points  
#correctly predicted/Total points predicted to be fraudulent  
count1 = 0  
total1 = 0  
for i in range(283726):  
    if clusters1[i]==0:  
        total1+=1;  
for i in range(283726):  
    if dataN.values[:,30][i]==1 and clusters1[i]==0:  
        count1+=1  
Precision1 = (count1/total1)*100  
Precision1
```

```
Out[16]: 0.11954254039319596
```

```
In [17]: #Find actual no. of fraudulent points in dataset for finding Recall  
data[data.Class==1].shape
```

```
Out[17]: (473, 31)
```

```
In [18]: #Recall of Fraudulent data Classification  
#Calculated as Recall = Number of fraudulent points  
#correctly predicted/Actual Number of fraudulent points  
count1 = 0  
for i in range(283726):  
    if dataN.values[:,30][i]==1 and clusters1[i]==0:  
        count1+=1  
Recall1 = (count1/473)*100  
Recall1
```

```
Out[18]: 33.192389006342495
```

```

In [19]: #Dropping those attributes which are not very correlated
         #with Class...on the basis of output obtained in Out[8]
traindata2 = traindata2.drop(labels=['V25', 'V15', 'V13', 'V26', 'V22', 'Amount', 'V23', \
                                     'V24', 'V28', 'Time', 'V20', 'V27', 'V21', 'V8', \
                                     'V19', 'V6', 'Class'],axis=1)

traindata2.shape

Out[19]: (283726, 14)

In [20]: #Performing Min_Max Normalization on traindata2
min_max_scaler2 = preprocessing.MinMaxScaler()
np_scaled2 = min_max_scaler2.fit_transform(traindata2)
traindata2N = pd.DataFrame(np_scaled2)
traindata2N.head()

Out[20]:
```


	0	1	2	3	4	5	6	7	8	\
0	0.93519	0.76649	0.88136	0.31302	0.76344	0.26682	0.47531	0.51060	0.25248	
1	0.97854	0.77007	0.84030	0.27180	0.76612	0.26488	0.45398	0.50527	0.38119	
2	0.93522	0.75312	0.86814	0.26877	0.76233	0.27018	0.41060	0.51302	0.32242	
3	0.94188	0.76530	0.86848	0.21366	0.76565	0.26680	0.41500	0.50759	0.27182	
4	0.93862	0.77652	0.86425	0.26980	0.76298	0.26897	0.49095	0.52430	0.23635	

	9	10	11	12	13
0	0.68091	0.63559	0.43439	0.73717	0.65507
1	0.74434	0.64122	0.46411	0.72779	0.64068
2	0.70668	0.64047	0.35744	0.76338	0.64495
3	0.71091	0.63637	0.41565	0.71125	0.78849
4	0.72448	0.60841	0.43500	0.72424	0.65067


```

In [21]: #Elbow method to find k for traindata2N
wcss2 = []
for i in range(1,20):
    kmeans2 = KMeans(n_clusters = i,init = 'k-means++',random_state = 0)
    kmeans2.fit(traindata2N)
    wcss2.append(kmeans2.inertia_)
plt.plot(range(1,20),wcss2)
plt.title('Fig3. Elbow Method - traindata2N')
plt.xlabel('Number of clusters')
plt.ylabel('WCSS2')
plt.show()

```



```
In [22]: #K-means clustering for traindata2N with k=8
kmeans2 = KMeans(n_clusters = 8,init = 'k-means++',random_state=0)
clusters2 = kmeans2.fit_predict(traindata2N)
clusters2
```

```
Out[22]: array([5, 6, 6, ..., 0, 5, 5], dtype=int32)
```

```
In [23]: #Accuracy of Classification.
#Consider cluster labeled by '7' to be fraudulent and
#that by all other labels to be Non-fraudulent
count2 = 0
for i in range(283726):
    if dataN.values[:,30][i]==0 and (clusters2[i]==0 or clusters2[i]==1 or \
                                     clusters2[i]==2 or clusters2[i]==3 or \
                                     clusters2[i]==4 or clusters2[i]==5 or \
                                     clusters2[i]==6):
        count2+=1
    if dataN.values[:,30][i]==1 and clusters2[i]==7:
        count2+=1
Accuracy2 = (count2/283726)*100
Accuracy2
```

```
Out[23]: 99.90977210407223
```

```

In [24]: #Precision of Fraudulent data classification
#Calculated as Precision = Number of fraudulent points
#correctly predicted/Total points predicted to be fraudulent
count2 = 0
total2 = 0
for i in range(283726):
    if clusters2[i]==7:
        total2+=1;
for i in range(283726):
    if dataN.values[:,30][i]==1 and clusters2[i]==7:
        count2+=1
Precision2_fraud = (count2/total2)*100
Precision2_fraud

```

Out[24]: 85.8085808580858

```

In [25]: #Precision of Non-Fraudulent data classification
#Calculated as Precision = Number of Non-fraudulent points
#correctly predicted/Total points predicted to be Non-fraudulent
count2 = 0
total2 = 0
for i in range(283726):
    if clusters2[i]==0 or clusters2[i]==1 or clusters2[i]==2 or clusters2[i]==3 or \
clusters2[i]==4 or clusters2[i]==5 or clusters2[i]==6:
        total2+=1;
for i in range(283726):
    if dataN.values[:,30][i]==0 and (clusters2[i]==0 or clusters2[i]==1 or \
clusters2[i]==2 or clusters2[i]==3 or \
clusters2[i]==4 or clusters2[i]==5 or \
clusters2[i]==6):
        count2+=1
Precision2_nonfraud = (count2/total2)*100
Precision2_nonfraud

```

Out[25]: 99.92484731302682

```

In [26]: #Recall of Fraudulent data Classification
#Calculated as Recall = Number of fraudulent points
#correctly predicted/Actual Number of fraudulent points
count2 = 0
for i in range(283726):
    if dataN.values[:,30][i]==1 and clusters2[i]==7:
        count2+=1
Recall2_fraud = (count2/473)*100
Recall2_fraud

```

Out[26]: 54.96828752642706

```

In [27]: #Recall of Non-Fraudulent data Classification
#Calculated as Recall = Number of Non-fraudulent points

```

```

#correctly predicted/Actual Number of Non-fraudulent points
count2 = 0
for i in range(283726):
    if dataN.values[:,30][i]==0 and (clusters2[i]==0 or clusters2[i]==1 or \
                                     clusters2[i]==2 or clusters2[i]==3 or \
                                     clusters2[i]==4 or clusters2[i]==5 or \
                                     clusters2[i]==6):
        count2+=1
Recall2_nonfraud = (count2/(283726-473))*100
Recall2_nonfraud

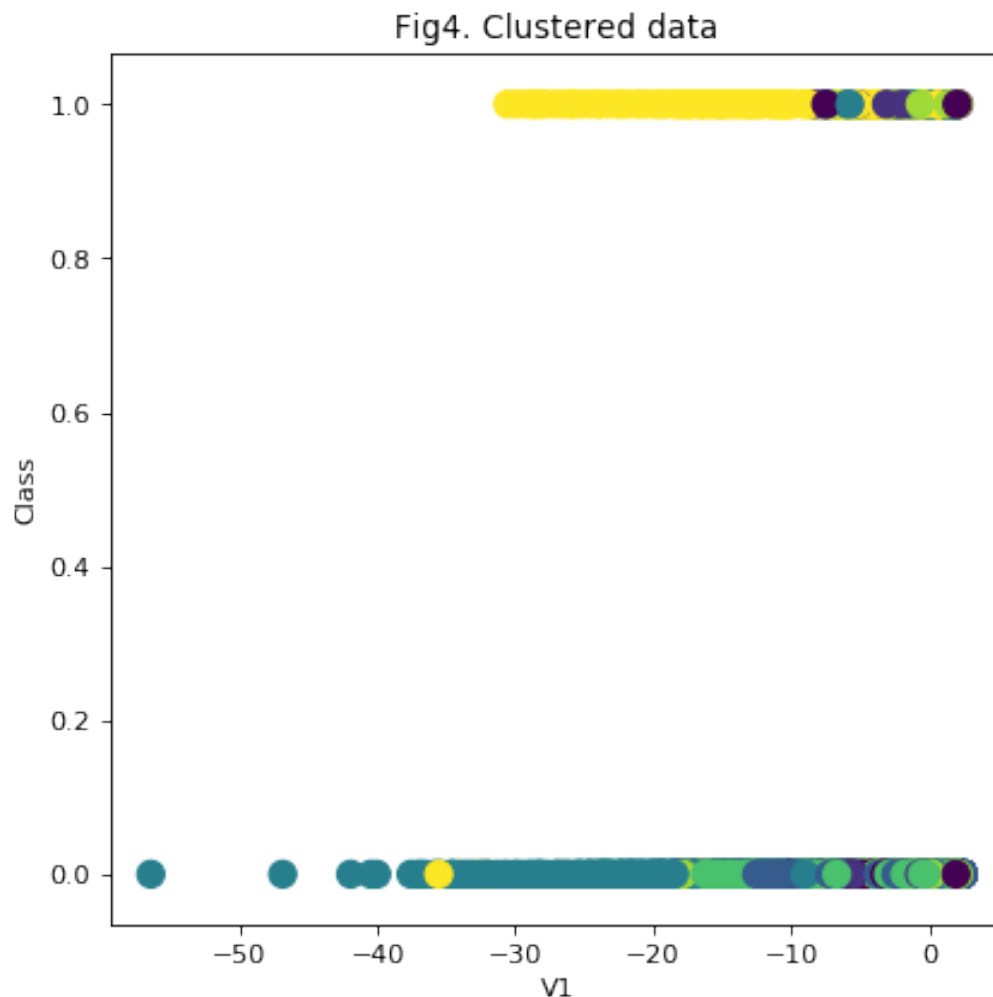
```

Out[27]: 99.98481922521562

```

In [28]: %matplotlib inline
plt.figure(figsize=(6, 6), dpi=80)
plt.scatter(data.iloc[:,1],data.iloc[:,30], c=clusters2,s=100, label='')
plt.title('Fig4. Clustered data')
plt.xlabel('V1')
plt.ylabel('Class')
plt.show()

```



```
In [29]: prediction = np.zeros(283726)
         for i in range(283726):
             if clusters2[i]==0 or clusters2[i]==1 or clusters2[i]==2 or clusters2[i]==3 or \
               clusters2[i]==4 or clusters2[i]==5 or clusters2[i]==6:
                 prediction[i]=0;
             if clusters2[i]==7:
                 prediction[i]=1;
         prediction
```

```
Out[29]: array([0., 0., 0., ..., 0., 0., 0.])
```

```
In [30]: pred = pd.DataFrame(prediction)
         pred.head()
```

```
Out[30]:      0
0 0.00000
1 0.00000
2 0.00000
3 0.00000
4 0.00000
```

```
In [31]: #Result
         Result = pd.concat([dataN.iloc[:,30],pred], axis=1)
         Result.columns = ['Actual Class','Predicted Class']
         Result
```

```
Out[31]:      Actual Class  Predicted Class
0          0.00000      0.00000
1          0.00000      0.00000
2          0.00000      0.00000
3          0.00000      0.00000
4          0.00000      0.00000
5          0.00000      0.00000
6          0.00000      0.00000
7          0.00000      0.00000
8          0.00000      0.00000
9          0.00000      0.00000
10         0.00000      0.00000
11         0.00000      0.00000
12         0.00000      0.00000
13         0.00000      0.00000
14         0.00000      0.00000
15         0.00000      0.00000
16         0.00000      0.00000
17         0.00000      0.00000
18         0.00000      0.00000
```

19	0.00000	0.00000
20	0.00000	0.00000
21	0.00000	0.00000
22	0.00000	0.00000
23	0.00000	0.00000
24	0.00000	0.00000
25	0.00000	0.00000
26	0.00000	0.00000
27	0.00000	0.00000
28	0.00000	0.00000
29	0.00000	0.00000
...
283696	0.00000	0.00000
283697	0.00000	0.00000
283698	0.00000	0.00000
283699	0.00000	0.00000
283700	0.00000	0.00000
283701	0.00000	0.00000
283702	0.00000	0.00000
283703	0.00000	0.00000
283704	0.00000	0.00000
283705	0.00000	0.00000
283706	0.00000	0.00000
283707	0.00000	0.00000
283708	0.00000	0.00000
283709	0.00000	0.00000
283710	0.00000	0.00000
283711	0.00000	0.00000
283712	0.00000	0.00000
283713	0.00000	0.00000
283714	0.00000	0.00000
283715	0.00000	0.00000
283716	0.00000	0.00000
283717	0.00000	0.00000
283718	0.00000	0.00000
283719	0.00000	0.00000
283720	0.00000	0.00000
283721	0.00000	0.00000
283722	0.00000	0.00000
283723	0.00000	0.00000
283724	0.00000	0.00000
283725	0.00000	0.00000

[283726 rows x 2 columns]

In [32]: *#Correlation between Predicted Class and Actual Class*
Result.corr()

Out [32]: Actual Class Predicted Class

Actual Class	1.00000	0.68639
Predicted Class	0.68639	1.00000

```
In [33]: #Root Mean Square Error  
rms = sqrt(mean_squared_error(Result.values[:,0],Result.values[:,1]))  
rms
```

```
Out[33]: 0.030037958640320633
```